

# Data Serialization

Paulo Sérgio Almeida

Grupo de Sistemas Distribuídos  
Departamento de Informática  
Universidade do Minho



# Kinds of heterogeneity

Heterogeneity in distributed systems may exist at several levels:

- Hardware;
- Programming Languages;
- Middleware.



# Hardware heterogeneity

Many reasons to mix different hardware:

- application availability for the OS/Hardware combination;
- gradual machine replacement over time;

Different hardware may mean:

- different data representation: need for neutral representation;
- different machine code: obstacle to code migration;
  - implementation for each machine;
  - use of virtual machines; compile once and run everywhere.



# Language heterogeneity

Many reasons to use/mix different languages:

- language suitability to problem;
- availability of libraries;
- legacy applications;

Tools to cope with language heterogeneity:

- data serialization formats + libraries;
- language neutral IDL + language mappings;
- code generation by IDL compiler; e.g. *stubs*, *skeletons*.



# Middleware heterogeneity

## Motivation:

- middleware *A* is popular; later middleware *B* is preferred;
- or, middleware *A* and *B* are suitable to different scenarios;
- two applications, built using *A* and *B*, need to interoperate;

## Multiple vendors of single middleware:

- different vendor implementations should interoperate;
- does not happen without a wire-protocol standard;
- e.g., CORBA before IIOP, with only language-level standards;

## Different middleware interoperability:

- interoperability protocols through *middleware bridges*;
- difficult for undefined/opaque/proprietary wire protocols;

Better to use simple middleware with defined wire-protocols.



# Data serialization and remote invocation protocols

- Data serialization formats allow data structures to be sent over the network; e.g. ONC XDR, CORBA CDR, JSON.
- Remote invocation protocols, defined over transport protocols such as TCP, are also defined; e.g., ONC RPC, CORBA IIOP, JSON-RPC.
- Commonly described over an underlying bidirectional *stream*;
- Ideally, we should be able to choose each layer independently;
- Many RPC frameworks bundle them together;



# Data serialization features

- Language specific versus language agnostic;  
e.g., Kryo versus JSON, CDR, Avro, Thrift, Protocol Buffers;
- Binary versus text-based (human readable);  
e.g., Kryo, CDR, Thrift, Protocol Buffers versus JSON;
- Does it allow sub-structure sharing / references?  
e.g., Kryo, CDR, YAML versus JSON, Avro, Thrift, Protocol Buffers;
- Versioning and Extensibility;  
does it allow different versions to co-exist?
- Self-description;  
how much data describes itself / needs external information;  
e.g. no info (e.g, XDR, CDR), field+type info (e.g., JSON);
- Encoding/decoding step vs neutral in-memory representation  
(binary formats);



# Versioning and Extensibility

- Versioning: to allow backward and forward compatibility along a sequence of versions;
  - adding field
  - renaming field
  - removing field
  - changing field type
- Distributed extensibility: to allow extensions concurrently developed to interoperate;
  - the same concerns as for versioning and
  - possible name collisions in new field
  - sometimes partially addressed with namespaces, which themselves introduce complexity;
  - complex problem; not our focus;





# Self-description

Where is the information that allows data to be decoded?

- Receiver code;
- External schema;
- Data stream itself;

Possible information in the data stream:

- Names of fields;
- Type information;



# Self-description: some possibilities (1)

- No description + code:
  - Sender and receiver (code) know exact structure;
  - Manual coding or IDL + code generation;
  - e.g., XDR, CORBA IIOP;
- No description + separate schema:
  - separate schema is used to decode the data;
  - e.g., Avro;
- Field tagging + type information + code;
  - sender and receiver (code) know expected structure;
  - extract relevant fields, skip unknown fields, default missing;
  - e.g., Thrift, Protocol Buffers;



## Self-description: some possibilities (2)

- Type and field names + separate schema for type information;
  - external schema contains type information to validate and decode;
  - e.g., some uses of XML;
- Field names + type information;
  - self-description of both field and type information;
  - e.g., JSON;



# Binary formats

- Examples: XDR, CORBA CDR, Thrift, Protocol Buffers;
- Compact and allow a computationally efficient serialization;
- May use standards (e.g., IEEE 754) for floatint point;
- Have many design options, e.g.:
  - Alignment in serialization stream;
  - Endianness (little/big endian) treatment;
  - Fixed vs variable size fields;



# Binary formats options/features (1)

- Alignment in serialization stream:
  - with alignment, resulting in *padding*; e.g., CDR;
  - without padding between fields; more compact; most formats;
- Endianness (little/big endian) treatment:
  - fixed neutral format; e.g., *network order*; most formats;
  - annotating format at sender; receiver converts only if different (*receiver makes it right*); more efficient; e.g. CDR.
- Fixed vs variable size fields:
  - For integer fields or describing the size of strings;
  - Most strings are small; most integers are small;
  - Fixed size assumes worst case;  
e.g., 4 byte string length, 8 byte integers in Thrift BinaryProtocol;
  - Variable size may use 1 or 2 bytes in the more common case; e.g., Protocol Buffers;



## Binary formats options/features (2)

- Outer object before vs encompassing inner objects:
  - possibility of incremental reads and processing partial message;
- Random access:
  - possibility of accessing one field without parsing the whole message
- Zero-copy:
  - possibility of accessing message in-place;
- Whether default fields take space on message;
- Whether pointers take space on message;
- Whether (de-serialized) message is usable as mutable state;
- Arena allocation vs individual allocation of each object;



# Neutral in-memory representation

- Most schemes involve an encoding/decoding step;
- Some binary formats use a neutral in-memory representation:
  - Cap'n Proto
  - Simple Binary Encoding
  - FlatBuffers
- Data is kept in memory already “serialized”;
  - neutral format, defining endianness, alignment, etc;
  - position independent through offset based pointers;
- Designed for random access and incremental access;
- Allows immediate sending or reading (e.g. mmaped file);
- Message usable only as immutable state;
- Allows inter-process communication by sharing memory;
  - even across different languages;



# Text based formats

- Convert binary values (e.g., numbers) into a textual representation;
- Less compact than binary formats;
- Much slower than binary formats (10 to 100 times);
  - conversion between binary and decimal number representation;
  - *parsing*;
- Increased popularity after WWW appeared;
- XML, popular with structured documents, was overused everywhere, including for data serialization;
- Other formats were proposed specifically for data serialization; e.g, YAML, JSON;
- JSON (JavaScript Object Notation) is becoming the more important text-based data serialization format;





# Problems with XML for data serialization

XML is appropriate in the *document centric* use, but very inappropriate in the *data centric* use:

- It does not support natively essential data structures, like lists/sequences/arrays and maps/dictionaries/hashes;
- Mappings from/to language data structures may be:
  - complex;
  - arbitrary (with no single natural normal form);
  - possibly incompatible between themselves;
- XML processing is complex, e.g.:
  - *callbacks* (SAX);
  - tree manipulation (DOM).
- Serialization is costly in size and very costly in processing time;



# Problems with XML for data serialization: arrays

- Lists/sequences/arrays are a fundamental data structure;
- With no native support in XML, they are mapped using one nested element for each array element;
- Some arbitrary element name has to be used;
- Corresponding opening and closing tags per array element results in much wasted space and parsing overhead;

```
<welcome-file-list>  
  <welcome-file>index.html</welcome-file>  
  <welcome-file>index.htm</welcome-file>  
  <welcome-file>index.jsp</welcome-file>  
</welcome-file-list>
```



# Problems with XML for data serialization: maps

- Maps/dictionaries/hashes are becoming important data structures, namely in dynamic languages.
- Maps are sets of key-value pairs;
- Problems in XML:
  - attributes are very limited, as they cannot hold arbitrary values;
  - this leads to using nested elements (but these are ordered);
- Key-value pairs are mapped to pairs of nested elements:

```
<member>  
  <key>ai</key>  
  <value>application/postscript</value>  
</member>
```



# YAML and JSON

- XML weakness for data serialization motivated alternative proposals (e.g., ConciseXML);
- Dynamic languages communities (Python, Perl, Ruby) defined YAML, a text based format appropriate for data serialization;
- With JavaScript, JSON became a popular text-based format;
- YAML 1.2 was standardized as a superset of JSON;
- Some YAML design goals were:
  - human friendly;
  - good interaction with dynamic languages;
  - native support for the more common data structures;
  - allow stream based processing;



# YAML/JSON vs XML

## Goals:

- XML descends from SGML, with emphasis on representing structured documents;
- YAML/JSON were designed from scratch for data serialization;

## Information model:

- an XML document is essentially a tree tagged with attributes;
  - each element contains an unordered list of attribute-value pairs;
  - and an ordered list of child elements;
- YAML/JSON uses data structure constructors: arrays, maps and scalars (e.g., numbers, strings, booleans),



# YAML/JSON: arrays

## Native array support in YAML/JSON:

- YAML:

- index.html
- index.htm
- index.jsp

- YAML/JSON:

```
["index.html", "index.htm", "index.jsp"]
```



# YAML/JSON: maps

## Native array support in YAML/JSON:

- **YAML:**

```
abs: audio/x-mpeg  
aif: audio/x-aiff  
aifc: audio/x-aiff  
aiff: audio/x-aiff
```

- **YAML/JSON:**

```
{ "abs": "audio/x-mpeg", "aif": "audio/x-aiff",  
  "aifc": "audio/x-aiff", "aiff": "audio/x-aiff" }
```



# YAML/JSON vs XML: arrays

## YAML:

```
welcome-file-list:  
  - index.html  
  - index.htm  
  - index.jsp
```

## YAML/JSON:

```
{ "welcome-file-list":  
  ["index.html",  
   "index.htm",  
   "index.jsp"] }
```

## XML:

```
<welcome-file-list>  
  <file>index.html</file>  
  <file>index.htm</file>  
  <file>index.jsp</file>  
</welcome-file-list>
```





# YAML/JSON vs XML: maps

## YAML/JSON:

```
{mime-mappings:  
  {"abs": "audio/x-mpeg",  
    "aif": "audio/x-aiff",  
    "aifc": "audio/x-aiff",  
    "aiff": "audio/x-aiff"}}}
```

## XML:

```
<mime-mapping>  
  <key>abs</key>  
  <value>audio/x-mpeg</value>  
</mime-mapping>  
<mime-mapping>  
  <key>aif</key>  
  <value>audio/x-aiff</value>  
</mime-mapping>  
<mime-mapping>  
  <key>aifc</key>  
  <value>audio/x-aiff</value>  
</mime-mapping>  
<mime-mapping>  
  <key>aiff</key>  
  <value>audio/x-aiff</value>  
</mime-mapping>
```



# Filters and filter composition

- A filter is an operation to write/read a data type to/from a stream;
- For composite types, invoke filter for each component;
- XDR filters are bidirectional: single function works both ways;

```
class A implements Streamable{  
    int i;  
    float f;  
    void readFrom(ObjInputStream str) {  
        this.i = str.readInt();  
        this.f = str.readFloat();  
    }  
    void writeTo(ObjOutputStream str) {  
        str.writeInt(this.i);  
        str.writeFloat(this.f);  
    }  
}
```



# Recursive structures


- To serialize recursive structures (e.g. lists, trees) references must be handled;
- Can be done writing boolean, followed by pointed structure serialization if non-null;
- A generic reference processing operation can be defined;

```
public class ObjOutputStream extends DataOutputStream {  
    public void writePtr(Streamable v) throws IOException {  
        if (v != null) {  
            this.writeBoolean(true);  
            v.writeTo(this);  
        } else  
            this.writeBoolean(false);  
    }  
  
    public class ObjInputStream extends DataInputStream {  
        public Streamable readPtr(Class c) throws IOException {  
            if (this.readBoolean())  
                return read(c);  
            else  
                return null;  
        }  
    }  
}
```



# Object graphs with sub-structure sharing and cycles

- Previous strategy does not handle sharing and cycles;
- Can be handled encoding index to previously encoded objects;
- Must keep map of encoded/decoded objects;



# Polymorphism: encoding type information

- If we know type of objects (fixed class), no need to encode;
- To allow subtype polymorphism, need to encode concrete type (e.g., class name in a string);
- Allows receiver to instantiate object and ask it to read itself;

```
public class ObjInputStream extends DataInputStream {  
    // ...  
    public Streamable readPoly() throws IOException {  
        String s = this.readUTF();  
        // ...  
        v = (Streamable)Class.forName(s).newInstance();  
        // ...  
        v.readFrom(this);  
        return v;  
    }  
}
```

