

# REST Architectural Style

Paulo Sérgio Almeida

Grupo de Sistemas Distribuídos  
Departamento de Informática  
Universidade do Minho



# Communication Architecture

- Saying that we have, e.g., distributed objects, does not tell much.
- It describes a programming paradigm oriented mostly towards remote invocation.
- Says nothing about actual communication architecture:
  - a few monolithic distributed objects? many small objects?
  - chatty RPC's with lots of ping-pong or bulky RPC's?
  - do we maintain session-state or are servers stateless?
  - do objects form a broker network to relay messages?
  - do we use any form of caching by intermediate proxies?



# Architectural styles in distributed systems

- Communication patterns between distributed components define architectural styles of distribution.
- An architectural style is defined as a set of constraints (of what is allowed to happen).
- Roy Fielding (one of the HTTP/1.1 authors) wrote the dissertation *Architectural Styles and the Design of Network-based Software Architectures*.
- Defines a taxonomy of architectural styles:
  - data-flow styles;
  - replication styles;
  - hierarchical styles;
  - mobile code styles;



# Data-flow styles

**Pipe and Filter (PF):** components form a *pipeline*, each acting as a filter; each reads a stream of data, produces stream of data; no state sharing between components;

**Uniform Pipe and Filter (UPF):** special case of PF; added constraint that all filters have the same interface; example: classic Unix pipes over byte streams;



# Replication styles

- Replication improves availability and scalability;
- Examples: *domain name system* (DNS), USENET news (nnntp);

**replicated repository (RR)**: several servers provide same service, providing illusion that there is only one; example: distributed file systems, remote versioning systems;

**cache (\$)**: variant in which result of individual request is kept to be reused by later requests, possibly from other clients; important for scalability, e.g., of WWW;



# Hierarchical styles

- client-server (CS): basic model from which others derive;
- layered-client-server (LCS): a server is client of other servers; introduces *overhead* and latency; allows intermediaries like *gateways* and *proxies*; style of classic n-tier systems;
- client-stateless-server (CSS): special case of CS in which there is no session state about clients in servers;
- client-cache-stateless-server (C\$SS): combines \$ with CSS; improves performance; example: Sun NFS;
- layered-client-cache-stateless-server (LC\$SS): example: DNS;
- remote session (RS): client initiates session in server, invokes commands and exits; state is kept on server;
- remote data access (RDA): client invokes series of operations that manipulate state on server or return parts of state; e.g., SQL commands;



# Mobile code

Brings processing closer to data, improving efficiency and remote resource utilization.

**virtual machine (VM):** same execution environment in each node; adds security and management problems; examples: applets in JVM, scripts in Javascript, PERL, Python.

**remote evaluation (REV):** combines *client-server* and *virtual machine*; client passes code and data to server, which performs computation and returns result;

**code on demand (COD):** client ask server for code and executes it locally; facilitates extensibility/upgrading of clients;

**layered-code-on-demand-client-cache-stateless-server (LCODC\$SS):** example: Web browser that allows applet downloading and protocol extensions;

**mobile agent (MA):** derivation of REV and COD, where an entire component, with code and data is moved to another site, possibly in the middle of a computation;



# REST: the architectural style of the WWW

- The most successful distributed service has been the WWW.
- Roy Fielding characterizes WWW architecture as REST.
- REST: Representational State Transfer
  - uniform interface, with a fixed set of operations;
  - resources and uniform resource identifiers (URIs);
  - hypermedia as the engine of application state (HATEOAS);
  - *layered client cache stateless server (LC\$SS)* style, with optional *code on demand (COD)*;
- These architectural constraints allowed extreme scalability, being important towards success of WWW.





# REST: uniform interface

- In REST, all components present same interface, with fixed set of operations, with universal semantics, regardless of application;
- E.g., in HTTP we have: GET, PUT, POST, DELETE, ...
- Contrasts with RPC/RMI systems, where different interfaces keep being defined for different distributed objects that may co-exist;
- Fixed interface important for long term evolution of large scale independently managed systems without global coordination;
- Universal semantics allows using generic intermediaries (like *proxies* and *gateways* in WWW), important for scalability;



# REST: resources and uniform resource identifiers

- REST is resource-oriented: a distributed system is organized as a set of resources;
- Each resource has a unique identifier (URI), in a single global name-space;
- All resources are accessed in a uniform way, through the same fixed interface;
- An application is designed in terms of resource design (definition, relationships); not in terms of interface definition;



# REST: hypermedia as the engine of application state (HATEOAS)

- Operations manipulate resources through *representations*;
- Example: an HTTP GET on a resource *transfers* a representation of that resource;
- Representations are manipulated through preferably standardized *media-types*;
- A representation may contain *links* to other resources;
- A client can progress by using one of the links in an obtained representation to invoke an operation on the respective resource;
- Hypermedia defines possible state transitions of clients;



# REST: layered client cache stateless server (LC\$SS)

- Interactions are stateless, with no conversational state:
  - a request must contain all information to be understood;
  - a component responds to a request without having to know previous requests from the client;
  - session state is kept in the client;
- Server state is only resource-related, not client-session related, leading to greater scalability and fault-tolerance;
- Intermediaries can understand each request separately;
- Clients or intermediaries can cache state when appropriate, improving scalability;



# Non-RESTful Web Services

- Distributed objects middleware (DCOM, CORBA, RMI) was not successful for interaction between independent organizations;
- Given the success of WWW, standards were defined using XML and working over HTTP: the *Web Services* (WS-\*);
- Those standards focus over aspects such as:
  - remote invocation (SOAP);
  - service definition: Web Services Description Language (WSDL);
  - discovery: Universal Description Discovery and Integration (UDDI);



# Non-RESTful Web Services critique

In the beginning of 2000s there was considerable polemic over WS-\*:

- WWW characteristics, present in REST architectural style, that made it successful are not respected;
- Example: interface definition and RPC orientation;
- Instead of embracing WWW, WS-\* merely use HTTP as a transport protocol, being unusable from a plain web application;
- Example: tunneling all requests through POST, with huge XML marshalling overhead, and not exploiting safety or idempotency;
- The use of HTTP as a transport protocol is formally incorrect;

*“HTTP is not a Transport Protocol” [Roy Fielding]*



# WS-\* versus RESTful use of HTTP + XML/JSON

- Oponents of WS-\* defend that:
  - instead of using a stack isolated from the web (SOAP, WSDL, UDDI) and interface-oriented;
  - WWW integration should be promoted, through the use of HTTP + XML/JSON according to REST principles (resource orientation);
- WWW integration and correct HTTP use brings advantages:
  - reutilization of existing mechanisms such as intermediaries (e.g., proxies), authentication, or access control;
  - allowing standard web clients (e.g., browsers, libcurl);
  - exploiting HTTP semantics (safety, idempotency) for scalability;



# HTTP methods

## Main methods:

- GET** retrieves resource representation;
- PUT** updates resource or creates new resource;
- POST** appends to resource or creates subordinate resource;
- DELETE** removes resource;

## Other methods:

- PATCH** modifies an existing resource;
- HEAD** retrieves metainformation about resource;
- OPTIONS** retrieves communication options for resource;
- TRACE** traces request along proxies;





# HTTP safe and idempotent methods

- A *safe* method should not have a visible effect upon resources (modify resource representation or create/remove resources)
- Method GET should be safe; never modify resources;
- An *idempotent* method should have the same effect if invoked several times (ignoring concurrent invocations by others);
- Methods GET, PUT and DELETE should be idempotent;
- POST has a “create new or append” semantics, being neither safe nor idempotent;



# HTTP safe and idempotent methods

Method	Safe	Idempotent
GET	yes	yes
HEAD	yes	yes
OPTIONS	yes	yes
TRACE	yes	yes
PUT	no	yes
DELETE	no	yes
POST	no	no
PATCH	no	no



# HTTP Status-Codes

- HTTP defines status-codes to be returned, to represent outcome;
- Most common fall in four classes: 2XX, 3XX, 4XX, 5XX;
- 2XX means successfully received, understood and accepted;
- 3XX means redirection; repeat request to new URI;
- 4XX means client error (e.g., not found, authorized or allowed);
- 5XX means server error, possibly temporary;



# Some common status-codes

200 OK
201 Created
204 No Content
301 Moved Permanently
304 Not Modified
400 Bad Request
401 Unauthorized
404 Not Found
405 Method Not Allowed
409 Conflict
412 Precondition Failed
415 Unsupported Media Type
500 Internal Server Error
501 Not Implemented
503 Service Unavailable



# GET method

- Should simply retrieve representation;
- Should be safe and idempotent;
- Should not modify, create or remove resources;
- Can be *conditional GET* using headers like If-Modified-Since;
- Conditional GET reduces network usage, refreshing client/cache without transferring representation;
- May be a *partial GET* using Range header, allowing partial data transmission; e.g. to restart transmission after network error;



# PUT method

- Sends entity to update resource of given URI;
- Should be idempotent;
- May also be used to create new resource with a given client defined URI (returning “201 Created” in this case);
- Invalidates cache entry for URI that may exist; response is not cacheable;



# POST method

- Sends entity to append to resource or create subordinate;
- Used to update or create resource in a non-idempotent way;
- Can add information to existing resource; examples:
  - increment counter;
  - add item to list;
- Can create subordinate resource with server-defined URI:
  - create new thread in forum;
  - create reply comment as child of another in a forum thread;
  - create new record in database table with server-defined key;
- When creating new resource should:
  - return “201 Created”;
  - return URI of created resource in “Location” header;



# DELETE method

- Removes resource;
- Is naturally idempotent;
- Invalidates cache entry for URI; response is not cacheable;





# Pointers

- **Roy Fielding dissertation:**  
`www.ics.uci.edu/~fielding/pubs/dissertation/  
fielding_dissertation.pdf`
- **HTTP/1.1:**  
`www.w3.org/Protocols/rfc2616/rfc2616.html`
- **Webmachine:**  
`www.github.com/basho/webmachine`
- **Dropwizard:**  
`dropwizard.io`
- **Example: Amazon S3 REST API:**  
`docs.aws.amazon.com/AmazonS3/latest/API/  
APIRest.html`
- **Example: last.fm REST API:**  
`www.last.fm/api/rest`

